

What did real waiver bids tell us?

A replay of five production strategies against 239 anonymized bids from SeaMex 2026. Useful signal, but not enough stable evidence to crown a default.

Evidence window

Reconstructed decision weeks W2 and W3

Data through

2026-09-25T23:00:08.539Z

Scope

One 32-team league · \$500 FAAB

BEST IN THIS SAMPLE

Aggressive

\$26.4 winning-bid MAE across n=47. It also leads the non-token winning subset at \$38.5.

ROBUST RECOMMENDATION

None

Leader changes under fixed outlier rules: Aggressive / Max VORP / Max VORP. The sample winner is not stable.

PRODUCT DECISION

Hold

Keep production unchanged. Collect exact pre-waiver snapshots across materially more weeks, then repeat this preregistered comparison.

Read this in 30 seconds

- Aggressive best matched the 47 observed winning prices, but its advantage depends on a few extreme targets.
- VoRP had the lowest MAE across all 239 usable bids and 165 serious bids; Max VORP had the strongest rank signal tied with VoRP.
- All projection snapshots used here are reconstructed, not exact historical forecasts. There are only 2 usable decision weeks.
- Manager-adjusted forecastable-win MAE is \$22.9 on only n=27; it is a smaller prior-history-eligible subset and is not directly comparable with the n=47 raw-strategy results.

“Weekly” is not a formula.


The production registry has Max VORP, VoRP, Safe, Aggressive, and Weeks as Starter. No authoritative Weekly key, alias, or implementation exists, so this report does not fabricate one.

Price fit depends on the question.

Winning-bid error asks who best estimates the clearing price. All-bid error also includes failed claims and token bids. Lower MAE is better; bias shows direction.

EVIDENCE VIEW

Winning-bid error by strategy




n=47 usable winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Lower MAE is better. Aggressive is best in this observed sample—not a universal winner.

EVIDENCE VIEW

All-bid vs. serious-bid error




All n=239; serious/non-token n=165 • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

VoRP has the lowest raw MAE in both broader bid subsets.

EVIDENCE VIEW

Signed bias on winning bids




n=47 winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Negative = underpredicts; zero is ideal.

EVIDENCE VIEW

Rank signal and serious-cluster coverage



Rank: n=47 wins; coverage: n=19 competitive target clusters • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Higher is better; the two measures answer different questions.

Winning-bid scorecard

n=47; subset=usable winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	47	28.0	15.2-44.5	7.0	-25.4	0.71	42.6%	\$0-\$285	\$0-\$84
● VoRP	47	29.1	15.8-45.6	7.0	-26.5	0.71	42.6%	\$0-\$285	\$0-\$77
● Safe	47	32.9	21.4-46.0	11.0	-9.1	0.61	36.2%	\$0-\$285	\$0-\$105
● Aggressive	47	26.4	17.2-37.6	11.0	-5.6	0.62	40.4%	\$0-\$285	\$0-\$184
● Weeks as Starter	47	29.4	17.4-44.1	11.0	-21.5	0.62	34.0%	\$0-\$285	\$0-\$105

All and serious scorecard — all bids

n=239; subset=all formula-usable bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	239	25.2	17.2-32.7	16.0	-2.1	0.72	27.6%	\$0-\$285	\$0-\$84
● VoRP	239	24.4	16.4-32.1	15.0	-5.2	0.72	29.7%	\$0-\$285	\$0-\$77
● Safe	239	39.2	32.8-44.2	39.0	+18.6	0.56	16.7%	\$0-\$285	\$0-\$105
● Aggressive	239	48.3	35.6-59.6	38.0	+37.8	0.60	19.7%	\$0-\$285	\$0-\$184
● Weeks as Starter	239	29.3	21.9-35.5	23.0	+4.1	0.61	23.0%	\$0-\$285	\$0-\$105

All and serious scorecard — serious bids

n=165; subset=non-token bids > \$5 • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1


STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	165	31.0	20.9-39.5	19.0	-8.3	0.57	18.2%	\$6-\$285	\$0-\$84
● VoRP	165	30.0	19.8-38.9	17.0	-12.6	0.57	21.2%	\$6-\$285	\$0-\$77
● Safe	165	41.0	34.4-46.8	38.0	+11.2	0.42	15.2%	\$6-\$285	\$0-\$105
● Aggressive	165	54.3	39.0-67.0	38.0	+39.3	0.48	15.2%	\$6-\$285	\$0-\$184
● Weeks as Starter	165	34.0	24.9-41.4	25.0	-2.4	0.49	16.4%	\$6-\$285	\$0-\$105

The winner changes when extremes move.

These filters were declared before inspecting winners and applied independently. They are sensitivity checks, not permission to delete inconvenient evidence.

EVIDENCE VIEW

Token and outlier sensitivity



Winning subsets n=47/29/28/26/25 in displayed order • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Aggressive leads the sample; Max VORP leads MAD/IQR-filtered views. No robust winner.

Token threshold

Low intent means $\text{bid} \leq \max(\$1, 1\% \text{ of original FAAB}) = \5 . Tokens remain in the all-bid view and are separated here.

Three outlier lenses

Ratio-gap flagged 1 isolated top; MAD flagged 3; IQR flagged 4. Different robust rules produce different leaders.

Honest conclusion


Aggressive is the sample winner. There is **no robust winner**. Treat this as a hypothesis, not a default-setting result.

Identity lines expose the misses.

Each dot is an anonymized winning target. A perfect forecast lies on the dashed diagonal; below it underpredicts. Extreme observed bids stretch beyond every intrinsic strategy; Aggressive reaches higher than the others, but still underpredicts the largest win.

EVIDENCE VIEW


Actual vs. predicted: where errors happen



$n=47$ winning bids per strategy • reconstructed W2-W3 • anonymized targets

R^2 is diagnostic only here and can be negative because these intrinsic valuations were not fitted to bids.

Manager-adjusted walk-forward forecasts



n=133 all forecastable / n=27 wins • strict prior-batch subset • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Separate subset: 106 of 239 usable bids lacked sufficient prior-only manager evidence.

Manager-adjusted subset

Strict prior-batch walk-forward • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1 • smaller, non-comparable subset

FORECAST SUBSET	N	MAE	MEDIAN AE	BIAS	SPEARMAN P	R ²	WITHIN TOLERANCE
All forecastable bids	133	28.0	18.0	+12.1	0.67	-0.07	31.6%
Forecastable wins	27	22.9	6.0	-19.9	0.76	0.46	51.9%

Measures, assumptions, and guardrails.

All formulas are evaluated on dollars unless marked normalized. Predictions are intrinsic strategy values; manager forecasts are capped at reconstructed pre-bid FAAB.

MAE

$$\text{MAE} = (1/n) \sum |\text{predicted} - \text{actual}|$$

Average dollar miss. Easy to interpret; extreme bids still matter.

Median absolute error

$$\text{Median AE} = \text{median}(|\text{predicted} - \text{actual}|)$$

The typical dollar miss; less sensitive to a small number of extremes.

Signed bias

$$\text{Bias} = (1/n) \sum (\text{predicted} - \text{actual})$$

Negative means systematic underprediction; positive means overprediction.

Spearman rank correlation

$$\rho = \text{Pearson correlation}(\text{rank}(\text{actual}), \text{rank}(\text{predicted}))$$

Tests whether expensive targets rank above cheap ones. Ties receive average ranks.

R² diagnostic

$$R^2 = 1 - \text{SSE} / \text{SST}$$

$\text{SSE} = \sum (\text{actual} - \text{predicted})^2$; $\text{SST} = \sum (\text{actual} - \text{mean}(\text{actual}))^2$. R² can be negative when a formula is worse than predicting the sample mean. That is valid here: intrinsic valuations were not fitted regressions, so R² is often unsuitable as a headline ranking.

Normalized FAAB error

$$\text{mean}(|\text{predicted}/\text{FAAB} - \text{actual}/\text{FAAB}|)$$

Reported against original \$500 and reconstructed pre-bid manager FAAB. It compares budget share, not willingness.

Cluster-bootstrap interval

$$2,000 \text{ resamples of player} \times \text{processing-batch clusters}$$

Fixed seed 20260927; percentile 2.5%–97.5%. Keeps correlated win/loss claims together. It does not repair snapshot-history error.

Coverage

$$\text{predicted} \in [\text{lowest serious loss}, \text{winning bid}]$$

Computed on 19 competitive player/batch clusters. It asks whether a strategy falls inside observed serious bidding range.

Outlier rules

$$\text{ratio} \geq 2 \times + \$50 \text{ gap}; \text{MAD} > \text{median} + 3 \times 1.4826 \times \text{MAD}; \text{IQR} > \text{Q3} + 1.5 \times \text{IQR}$$

MAD/IQR also require a \$25 gap; minimum cluster sizes are 3 and 4. Flags are independent and never silently deleted.

Within tolerance

$$|\text{predicted} - \text{actual}| \leq \max(\$5, 20\% \text{ of actual})$$

A practical accuracy band shown alongside continuous error measures.

Provenance boundary

All 239 formula-usable events use same-week reconstructed W2/W3 projection snapshots and transaction-ledger FAAB reconstruction. Reconstructed means captured after the canonical cutoff; it is not proof of what managers saw. W4 was excluded because only an earlier-week fallback existed. No identities, league ID, player names/IDs, manager names/IDs, raw payload, or credentials appear in this artifact. Regenerate deterministically with `npm run analyze:bidding:presentation`.