

What did real waiver bids tell us?

A replay of five production strategies against 239 anonymized bids from SeaMex 2026. Useful signal, but not enough stable evidence to crown a default.

Evidence window

Reconstructed decision weeks W2 and W3

Data through

2026-09-25T23:00:08.539Z

Scope

One 32-team league · \$500 FAAB

Five-bullet answer

- Scope: W2 and W3 have same-week reconstructed snapshots. The top three are unique canonical winners after owner-directed-outlier-01 (private GET-only catalog match plus exact position/total fingerprints in both reconstructed snapshots map uniquely to one canonical completed W3 win at \$234; proof matches=1); the closest overall serious-median strategy is VoRP.
- W2 top-three winning-price arithmetic multipliers (observed/intrinsic): Max VORP 3.34× (3/3); VoRP 3.38× (3/3); Safe 2.95× (3/3); Weeks as Starter 3.37× (3/3).
- W3 top-three winning-price arithmetic multipliers (observed/intrinsic): Max VORP 2.35× (3/3); VoRP 3.00× (3/3); Safe 2.87× (3/3); Weeks as Starter 2.87× (3/3).
- W2 top-three serious-market-median arithmetic multipliers: Max VORP 1.40× (3/3); VoRP 1.42× (3/3); Safe 1.35× (3/3); Weeks as Starter 1.53× (3/3); closest across all W2 eligible market clusters: Max VORP.
- W3 top-three serious-market-median arithmetic multipliers: Max VORP 0.78× (3/3); VoRP 1.00× (3/3); Safe 1.13× (3/3); Weeks as Starter 1.13× (3/3); closest across all W3 eligible market clusters: VoRP.

Scope guardrail

Only W2 and W3 have like-for-like same-week reconstructed snapshots. W4 is explained as missing rather than backfilled. Aggressive is omitted here because it is derived from Safe.

CLOSEST MEDIAN MARKET

VoRP

Lowest MAE across owner-directed player/week clusters; competing claims stay within the canonical winner's processing batch.

STRATEGY × SCALE

Separate them

Intrinsic formulas describe target shape; observed/intrinsic multipliers estimate weekly market scale.

PRODUCT DECISION


Hold

Two reconstructed weeks do not support manager-style claims or a production default change.

Shape is not scale.

EVIDENCE VIEW

Weekly top-three winning multipliers




Owner-directed view (owner-directed-outlier-01 excluded); same-week snapshots only • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Arithmetic mean of winning/intrinsic ratios across each week's top-three unique winners.

Top-three canonical winners and serious market medians use observed bid ÷ intrinsic same-week suggestion. Arithmetic mean is primary; geometric mean, median, and coverage expose ratio skew and zero denominators.

EVIDENCE VIEW

Weekly top-three serious-market multipliers



Owner-directed view (owner-directed-outlier-01 excluded); same-week snapshots only • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Arithmetic mean of serious-median/intrinsic ratios across each week's top-three unique winners.

Exact weekly top-three multipliers

Observed/intrinsic; arithmetic mean is primary, with geometric mean, median, and zero-denominator coverage.

WEEK	OBSERVATION	STRATEGY	COVERAGE	ARITHMETIC MEAN	GEOMETRIC MEAN	MEDIAN
W2	Winning bid	Max VORP	3/3	3.34×	3.21×	2.73×
W2	Winning bid	VoRP	3/3	3.38×	3.24×	2.78×
W2	Winning bid	Safe	3/3	2.95×	2.78×	3.17×
W2	Winning bid	Weeks as Starter	3/3	3.37×	3.24×	3.65×
W2	Serious-market median	Max VORP	3/3	1.40×	1.40×	1.42×
W2	Serious-market median	VoRP	3/3	1.42×	1.42×	1.44×
W2	Serious-market median	Safe	3/3	1.35×	1.22×	0.94×
W2	Serious-market median	Weeks as Starter	3/3	1.53×	1.41×	1.09×
W3	Winning bid	Max VORP	3/3	2.35×	2.17×	2.23×
W3	Winning bid	VoRP	3/3	3.00×	2.79×	2.43×
W3	Winning bid	Safe	3/3	2.87×	2.73×	3.00×
W3	Winning bid	Weeks as Starter	3/3	2.87×	2.73×	3.00×
W3	Serious-market median	Max VORP	3/3	0.78×	0.78×	0.79×
W3	Serious-market median	VoRP	3/3	1.00×	1.00×	1.03×
W3	Serious-market median	Safe	3/3	1.13×	0.98×	0.73×
W3	Serious-market median	Weeks as Starter	3/3	1.13×	0.98×	0.73×

Top-three winning-bid ratios (privacy-safe labels)

Canonical winning bid ÷ intrinsic suggestion; exact FAAB equality is labeled censored.

TARGET	WEEK	WINNING BID	FAAB-CENSORED	MAX VORP	VORP	SAFE	WEEKS AS STARTER
W2 target 1	W2	\$285	no	4.75×	4.83×	3.17×	3.65×
W2 target 2	W2	\$153	no	2.73×	2.78×	1.72×	2.15×
W2 target 3	W2	\$99	no	2.54×	2.54×	3.96×	4.30×
W3 target 1	W3	\$231	no	3.50×	4.62×	3.00×	3.00×
W3 target 2	W3	\$187	no	2.23×	2.43×	1.78×	1.78×
W3 target 3	W3	\$103	no	1.32×	1.94×	3.81×	3.81×

Top-three serious-market ratios (privacy-safe labels)

Serious median ÷ intrinsic suggestion; competitors are canonical legitimate failures with bid > \$5.

TARGET	WEEK	SERIOUS MEDIAN	CENSORED OBS.	MAX VORP	VORP	SAFE	WEEKS AS STARTER
W2 target 1	W2	\$85.0	0	1.42×	1.44×	0.94×	1.09×
W2 target 2	W2	\$75.0	0	1.34×	1.36×	0.84×	1.06×
W2 target 3	W2	\$56.5	0	1.45×	1.45×	2.26×	2.46×
W3 target 1	W3	\$56.0	0	0.85×	1.12×	0.73×	0.73×
W3 target 2	W3	\$66.5	0	0.79×	0.86×	0.63×	0.63×
W3 target 3	W3	\$54.5	0	0.70×	1.03×	2.02×	2.02×

Median-market fit: overall, by week, and all-bid sensitivity

★ lowest MAE within scope; bias = intrinsic – observed. R² and Spearman appear only when meaningful.

SCOPE	OBSERVATION	STRATEGY	N	MAE	MEDIAN AE	BIAS	R ²	SPEARMAN P
Overall	Serious median	Max VORP	21	12.5	8.0	+1.5	0.56	0.79
Overall	Serious median	VoRP ★	21	11.0	7.0	-1.0	0.63	0.78
Overall	Serious median	Safe	25	28.4	31.5	+18.7	-1.26	0.42
Overall	Serious median	Weeks as Starter	23	17.4	12.0	+5.9	0.11	0.57
Overall	All-bid sensitivity	Max VORP	25	12.1	9.0	+6.4	0.66	0.86
Overall	All-bid sensitivity	VoRP ★	25	10.6	8.5	+4.3	0.71	0.86
Overall	All-bid sensitivity	Safe	34	28.2	25.0	+22.9	-1.10	0.50
Overall	All-bid sensitivity	Weeks as Starter	32	15.2	10.5	+9.2	0.31	0.69
W2	Serious median	Max VORP ★	11	14.9	14.0	-3.7	0.52	0.81
W2	Serious median	VoRP	11	15.0	14.0	-4.0	0.52	0.81
W2	Serious median	Safe	11	31.0	31.5	+25.0	-1.05	0.46
W2	Serious median	Weeks as Starter	11	18.7	12.0	+5.0	0.14	0.50
W2	All-bid sensitivity	Max VORP ★	14	12.6	10.0	+2.9	0.65	0.81
W2	All-bid sensitivity	VoRP	14	12.7	10.0	+2.6	0.65	0.81
W2	All-bid sensitivity	Safe	15	31.3	28.0	+27.6	-0.90	0.64
W2	All-bid sensitivity	Weeks as Starter	15	16.6	10.5	+9.3	0.33	0.64
W3	Serious median	Max VORP	10	9.9	8.0	+7.1	0.65	0.91
W3	Serious median	VoRP ★	10	6.6	6.0	+2.3	0.85	0.91
W3	Serious median	Safe	14	26.3	30.3	+13.8	-1.59	0.48
W3	Serious median	Weeks as Starter	12	16.3	12.8	+6.8	0.05	0.73
W3	All-bid sensitivity	Max VORP	11	11.4	8.5	+10.9	0.66	0.89
W3	All-bid sensitivity	VoRP ★	11	7.9	8.0	+6.5	0.83	0.89
W3	All-bid sensitivity	Safe	19	25.7	25.0	+19.3	-1.40	0.44
W3	All-bid sensitivity	Weeks as Starter	17	13.9	10.5	+9.1	0.27	0.73

Compact with-vs-without sensitivity

private GET-only catalog match plus exact position/total fingerprints in both reconstructed snapshots map uniquely to one canonical completed W3 win at \$234; marker only, with no private identity stored.

VIEW	CLUSTERS	CLOSEST	MAX VORP MAE	VORP MAE	SAFE MAE	WEEKS AS STARTER MAE
Owner-directed exclusion	28	VoRP	\$12.5	\$11.0	\$28.4	\$17.4
Raw evidence retained	29	VoRP	\$12.5	\$11.1	\$28.5	\$16.9

Interpretation boundary


The owner-directed marker is excluded only from this labeled analysis; raw canonical evidence remains in the sensitivity row. Serious means bid > \$5, failed claims must be legitimate same-target competitors, and FAAB-limited observations are marked censored. This supports a strategy-shape × market-scale hypothesis, not individual manager-style conclusions from two reconstructed weeks.

Price fit depends on the question.

Winning-bid error asks who best estimates the clearing price. All-bid error also includes failed claims and token bids. Lower MAE is better; bias shows direction.

EVIDENCE VIEW

Winning-bid error by strategy




n=47 usable winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Lower MAE is better. Aggressive is best in this observed sample—not a universal winner.

EVIDENCE VIEW

All-bid vs. serious-bid error




All n=239; serious/non-token n=165 • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

VoRP has the lowest raw MAE in both broader bid subsets.

EVIDENCE VIEW

Signed bias on winning bids




n=47 winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Negative = underpredicts; zero is ideal.

EVIDENCE VIEW

Rank signal and serious-cluster coverage



Rank: n=47 wins; coverage: n=19 competitive target clusters • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Higher is better; the two measures answer different questions.

Winning-bid scorecard

n=47; subset=usable winning bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	47	28.0	15.2-44.5	7.0	-25.4	0.71	42.6%	\$0-\$285	\$0-\$84
● VoRP	47	29.1	15.8-45.6	7.0	-26.5	0.71	42.6%	\$0-\$285	\$0-\$77
● Safe	47	32.9	21.4-46.0	11.0	-9.1	0.61	36.2%	\$0-\$285	\$0-\$105
● Aggressive	47	26.4	17.2-37.6	11.0	-5.6	0.62	40.4%	\$0-\$285	\$0-\$184
● Weeks as Starter	47	29.4	17.4-44.1	11.0	-21.5	0.62	34.0%	\$0-\$285	\$0-\$105

All and serious scorecard — all bids

n=239; subset=all formula-usable bids • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	239	25.2	17.2-32.7	16.0	-2.1	0.72	27.6%	\$0-\$285	\$0-\$84
● VoRP	239	24.4	16.4-32.1	15.0	-5.2	0.72	29.7%	\$0-\$285	\$0-\$77
● Safe	239	39.2	32.8-44.2	39.0	+18.6	0.56	16.7%	\$0-\$285	\$0-\$105
● Aggressive	239	48.3	35.6-59.6	38.0	+37.8	0.60	19.7%	\$0-\$285	\$0-\$184
● Weeks as Starter	239	29.3	21.9-35.5	23.0	+4.1	0.61	23.0%	\$0-\$285	\$0-\$105

All and serious scorecard — serious bids

n=165; subset=non-token bids > \$5 • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1


STRATEGY	N	MAE ↓	95% BOOTSTRAP MAE CI	MEDIAN AE ↓	BIAS	SPEARMAN P ↑	WITHIN TOLERANCE ↑	ACTUAL RANGE	PREDICTED RANGE
● Max VORP	165	31.0	20.9-39.5	19.0	-8.3	0.57	18.2%	\$6-\$285	\$0-\$84
● VoRP	165	30.0	19.8-38.9	17.0	-12.6	0.57	21.2%	\$6-\$285	\$0-\$77
● Safe	165	41.0	34.4-46.8	38.0	+11.2	0.42	15.2%	\$6-\$285	\$0-\$105
● Aggressive	165	54.3	39.0-67.0	38.0	+39.3	0.48	15.2%	\$6-\$285	\$0-\$184
● Weeks as Starter	165	34.0	24.9-41.4	25.0	-2.4	0.49	16.4%	\$6-\$285	\$0-\$105

The winner changes when extremes move.

These filters were declared before inspecting winners and applied independently. They are sensitivity checks, not permission to delete inconvenient evidence.

EVIDENCE VIEW

Token and outlier sensitivity



Winning subsets n=47/29/28/26/25 in displayed order • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Aggressive leads the sample; Max VORP leads MAD/IQR-filtered views. No robust winner.

Token threshold

Low intent means $\text{bid} \leq \max(\$1, 1\% \text{ of original FAAB}) = \5 . Tokens remain in the all-bid view and are separated here.

Three outlier lenses

Ratio-gap flagged 1 isolated top; MAD flagged 3; IQR flagged 4. Different robust rules produce different leaders.

Honest conclusion


Aggressive is the sample winner. There is **no robust winner**. Treat this as a hypothesis, not a default-setting result.

Identity lines expose the misses.

Each dot is an anonymized winning target. A perfect forecast lies on the dashed diagonal; below it underpredicts. Extreme observed bids stretch beyond every intrinsic strategy; Aggressive reaches higher than the others, but still underpredicts the largest win.

EVIDENCE VIEW


Actual vs. predicted: where errors happen



$n=47$ winning bids per strategy • reconstructed W2-W3 • anonymized targets

R^2 is diagnostic only here and can be negative because these intrinsic valuations were not fitted to bids.

Manager-adjusted walk-forward forecasts



n=133 all forecastable / n=27 wins • strict prior-batch subset • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1

Separate subset: 106 of 239 usable bids lacked sufficient prior-only manager evidence.

Manager-adjusted subset

Strict prior-batch walk-forward • Reconstructed W2-W3 projection snapshots • transaction-ledger FAAB • anonymized fixture v1 • smaller, non-comparable subset

FORECAST SUBSET	N	MAE	MEDIAN AE	BIAS	SPEARMAN P	R²	WITHIN TOLERANCE
All forecastable bids	133	28.0	18.0	+12.1	0.67	-0.07	31.6%
Forecastable wins	27	22.9	6.0	-19.9	0.76	0.46	51.9%

Measures, assumptions, and guardrails.

All formulas are evaluated on dollars unless marked normalized. Predictions are intrinsic strategy values; manager forecasts are capped at reconstructed pre-bid FAAB.

MAE

$$\text{MAE} = (1/n) \sum |\text{predicted} - \text{actual}|$$

Average dollar miss. Easy to interpret; extreme bids still matter.

Median absolute error

$$\text{Median AE} = \text{median}(|\text{predicted} - \text{actual}|)$$

The typical dollar miss; less sensitive to a small number of extremes.

Signed bias

$$\text{Bias} = (1/n) \sum (\text{predicted} - \text{actual})$$

Negative means systematic underprediction; positive means overprediction.

Spearman rank correlation

$$\rho = \text{Pearson correlation}(\text{rank}(\text{actual}), \text{rank}(\text{predicted}))$$

Tests whether expensive targets rank above cheap ones. Ties receive average ranks.

R² diagnostic

$$R^2 = 1 - \text{SSE} / \text{SST}$$

$\text{SSE} = \sum (\text{actual} - \text{predicted})^2$; $\text{SST} = \sum (\text{actual} - \text{mean}(\text{actual}))^2$. R² can be negative when a formula is worse than predicting the sample mean. That is valid here: intrinsic valuations were not fitted regressions, so R² is often unsuitable as a headline ranking.

Normalized FAAB error

$$\text{mean}(|\text{predicted}/\text{FAAB} - \text{actual}/\text{FAAB}|)$$

Reported against original \$500 and reconstructed pre-bid manager FAAB. It compares budget share, not willingness.

Cluster-bootstrap interval

$$2,000 \text{ resamples of player} \times \text{processing-batch clusters}$$

Fixed seed 20260927; percentile 2.5%–97.5%. Keeps correlated win/loss claims together. It does not repair snapshot-history error.

Coverage

$$\text{predicted} \in [\text{lowest serious loss}, \text{winning bid}]$$

Computed on 19 competitive player/batch clusters. It asks whether a strategy falls inside observed serious bidding range.

Outlier rules

$$\text{ratio} \geq 2 \times + \$50 \text{ gap}; \text{MAD} > \text{median} + 3 \times 1.4826 \times \text{MAD}; \text{IQR} > \text{Q3} + 1.5 \times \text{IQR}$$

MAD/IQR also require a \$25 gap; minimum cluster sizes are 3 and 4. Flags are independent and never silently deleted.

Within tolerance

$$|\text{predicted} - \text{actual}| \leq \max(\$5, 20\% \text{ of actual})$$

A practical accuracy band shown alongside continuous error measures.

Provenance boundary

All 239 formula-usable events use same-week reconstructed W2/W3 projection snapshots and transaction-ledger FAAB reconstruction. Reconstructed means captured after the canonical cutoff; it is not proof of what managers saw. W4 was excluded because only an earlier-week fallback existed. No identities, league identifier, player names/IDs, manager names/IDs, raw payload, or credentials appear in this artifact. Regenerate deterministically with `npm run analyze:bidding:presentation`.